

raw page (crop)

GANN GERTZ

3.3.2. ANSWER GENERATION

Vakharia et al. (2024) propose to employ an LLM to generate a final user-friendly answer based on the Prolog result and the initial user question. Additionally, Yang et al. (2025) propose to take advantage of the traces of Prolog executions as an explanation for the generated answers. In our approach, the Prolog trace additionally contains information that can uniquely map each rule and fact to a source chunk and the document it was extracted from. This enables further verification and attribution of reasoning steps.

4. Experiments

To evaluate the effectiveness of NeSy-RAG, we implemented a prototype and assessed it on the ShARC (Saeidi et al., 2018) benchmark, a conversational question answering dataset built from legal text sourced from the UK government website. Each entry consists of a **snippet** (a passage from the website), a **question**, a **scenario** describing the user’s situation, a **history** of prior follow-up exchanges, **evidence** capturing user-specific information extractable from the scenario, and a target **answer** (yes, no, irrelevant or more).

PRISM output, compiled

3.3.2. ANSWER GENERATION

GANN GERTZ

Vakharia et al. (2024) propose to employ an LLM to generate a final user-friendly answer based on the Prolog result and the initial user question. Additionally, Yang et al. (2025) propose to take advantage of the traces of Prolog executions as an explanation for the generated answers. In our approach, the Prolog trace additionally contains information that can uniquely map each rule and fact to a source chunk and the document it was extracted from. This enables further verification and attribution of reasoning steps.

4. Experiments

To evaluate the effectiveness of NeSy-RAG, we implemented a prototype and assessed it on the ShARC (Saeidi et al., 2018) benchmark, a conversational question answering dataset built from legal text sourced from the UK government website. Each entry consists of a snippet (a passage from the website), a question, a scenario describing the user’s situation, a history of prior follow-up exchanges, evidence capturing user-specific information extractable from the scenario, and a target answer (yes, no, irrelevant or more). Rather than supplying all user information upfront, we consolidate scenario, evidence, and history into a single context field that is withheld from the system at the start of each interaction. This design simulates a setting in which the system must actively recover missing user information through follow-up questions. Whenever the system poses a followup question, a separate LLM (as part of the evaluation system) answers it by consulting the context, then by simulating a follow-up question.